

Two-Way ANOVA Handout

1 Initialization

```
> library(NCStats)
> library(gdata)      # for aggregate.table()
> library(multcomp)    # for glht()
```

2 Bacteria Example

```
> bact <- read.table("Bacteria.txt",header=TRUE)
> str(bact)
```

```
'data.frame':      30 obs. of  3 variables:
 $ temp : int  27 27 27 27 27 35 35 35 35 35 ...
 $ conc  : num  0.6 0.8 1 1.2 1.4 0.6 0.8 1 1.2 1.4 ...
 $ cells: int  55 120 186 260 151 82 166 179 223 178 ...
```

```
> bact$ftemp <- factor(bact$temp)
> bact$fconc <- factor(bact$conc)
> str(bact)
```

```
'data.frame':      30 obs. of  5 variables:
 $ temp : int  27 27 27 27 27 35 35 35 35 35 ...
 $ conc  : num  0.6 0.8 1 1.2 1.4 0.6 0.8 1 1.2 1.4 ...
 $ cells: int  55 120 186 260 151 82 166 179 223 178 ...
 $ ftemp: Factor w/ 3 levels "27","35","43": 1 1 1 1 1 2 2 2 2 2 ...
 $ fconc: Factor w/ 5 levels "0.6","0.8","1",...: 1 2 3 4 5 1 2 3 4 5 ...
```

Explore sample size per group, univariate statistics, statistics for each group, and table of group means,

```
> with(bact, table(ftemp, fconc))
```

```
      fconc
ftemp 0.6 0.8 1 1.2 1.4
  27   2   2  2   2   2
  35   2   2  2   2   2
  43   2   2  2   2   2
```

```
> Summarize(cells~ftemp, data=bact)
```

```
  ftemp   n  Mean St. Dev. Min. 1st Qu. Median 3rd Qu. Max.
1     27  10 153.1 69.83544   55  113.2  141.5  177.2  274
2     35  10 169.3 51.14044   82  157.0  172.0  209.0  236
3     43  10 161.5 33.95176  115  138.0  149.5  185.2  228
```

```
> Summarize(cells~fconc, data=bact)
```

```

      fconc n      Mean St. Dev. Min. 1st Qu. Median 3rd Qu. Max.
1    0.6 6 108.1667 38.77843   55    85.0  104.5   141.2  153
2    0.8 6 144.1667 34.00245   91   126.5  151.0   163.5  186
3      1 6 154.8333 23.69318  133   135.2  148.0   174.2  186
4    1.2 6 235.0000 29.85297  189   224.2  232.0   254.0  274
5    1.4 6 164.3333 37.30773  111   145.8  164.5   181.8  219

```

```
> Summarize(cells~fconc, numdigs=2, data=bact)
```

```

      ftemp fconc n      Mean St. Dev. Min. 1st Qu. Median 3rd Qu. Max.
1      27    0.6 2 102.5    67.18   55   78.75  102.5  126.2  150
2      35    0.6 2  88.0     8.49   82   85.00   88.0   91.0   94
3      43    0.6 2 134.0    26.87  115  124.50  134.0  143.5  153
4      27    0.8 2 105.5    20.51   91   98.25  105.5  112.8  120
5      35    0.8 2 161.0     7.07  156  158.50  161.0  163.5  166
6      43    0.8 2 166.0    28.28  146  156.00  166.0  176.0  186
7      27      1 2 159.5    37.48  133  146.20  159.5  172.8  186
8      35      1 2 169.5    13.44  160  164.80  169.5  174.2  179
9      43      1 2 135.5     0.71  135  135.20  135.5  135.8  136
10     27     1.2 2 267.0     9.90  260  263.50  267.0  270.5  274
11     35     1.2 2 229.5     9.19  223  226.20  229.5  232.8  236
12     43     1.2 2 208.5    27.58  189  198.80  208.5  218.2  228
13     27     1.4 2 131.0    28.28  111  121.00  131.0  141.0  151
14     35     1.4 2 198.5    28.99  178  188.20  198.5  208.8  219
15     43     1.4 2 163.5    27.58  144  153.80  163.5  173.2  183

```

```
> with(bact, aggregate.table(cells, ftemp, fconc, FUN=mean))
```

```

      0.6  0.8    1   1.2   1.4
27 102.5 105.5 159.5 267.0 131.0
35  88.0 161.0 169.5 229.5 198.5
43 134.0 166.0 135.5 208.5 163.5

```

Fit the model and get the ANOVA results

```
> lm1 <- lm(cells~fconc, data=bact)
> anova(lm1)
```

Analysis of Variance Table

```



Response: cells
      Df Sum Sq Mean Sq F value    Pr(>F)
fconc   4  51596 12899.1  16.8154 2.041e-05
fconc:fconc 8  14703  1837.8   2.3958  0.06886
Residuals 15  11507    767.1

```

Examine the interaction plots



```
> fitPlot(lm1)
> fitPlot(lm1, interval=FALSE, change.order=TRUE, xlab="Concentration (%)",
  ylab="Mean Number of Cells", legend="topleft")

```



Examine the main effects plots

```
> fitPlot(lm1,which="ftemp",xlab="Temperature (C)",
  ylab="Mean Number of Cells",ylim=c(60,270))
> fitPlot(lm1,which="fconc",xlab="Concentration (%)",
  ylab="Mean Number of Cells",ylim=c(60,270))
```



Examine Tukey's HSD results for concentration main effect and construct a main effects plot with significance letters.

```
> bact.mc1 <- glht(lm1,mcp(fconc="Tukey"))
> summary(bact.mc1)
```

Simultaneous Tests for General Linear Hypotheses

Multiple Comparisons of Means: Tukey Contrasts


```
Fit: lm(formula = cells ~ ftemp * fconc, data = bact)
```

Linear Hypotheses:

	Estimate	Std. Error	t value	Pr(> t)
0.8 - 0.6 == 0	3.0	27.7	0.108	0.999965
1 - 0.6 == 0	57.0	27.7	2.058	0.287200
1.2 - 0.6 == 0	164.5	27.7	5.939	0.000196
1.4 - 0.6 == 0	28.5	27.7	1.029	0.838190
1 - 0.8 == 0	54.0	27.7	1.950	0.334998
1.2 - 0.8 == 0	161.5	27.7	5.831	0.000275
1.4 - 0.8 == 0	25.5	27.7	0.921	0.884508
1.2 - 1 == 0	107.5	27.7	3.881	0.010932
1.4 - 1 == 0	-28.5	27.7	-1.029	0.838184
1.4 - 1.2 == 0	-136.0	27.7	-4.910	0.001464

(Adjusted p values reported -- single-step method)

```
> fitPlot(lm1,which="fconc",xlab="Concentration (%)",ylab="Mean Number of Cells",main="")
> addSigLetters(lm1,which="fconc",lets=c("a","ab","ab","c","b"),pos=c(2,2,4,2,4))
```



3 Soil Phosphorous Example

You must change the directory to where the following file is located.

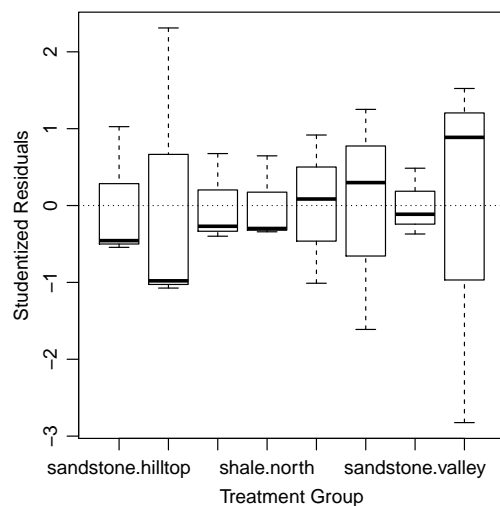
```
> sp <- read.table("SoilPhosphorous.txt",header=TRUE)
> str(sp)

'data.frame':      24 obs. of  3 variables:
 $ soil: Factor w/ 2 levels "sandstone","shale": 2 2 2 2 2 2 2 2 2 2 ...
 $ topo: Factor w/ 4 levels "hilltop","north",...: 4 4 4 2 2 2 3 3 3 1 ...
 $ phos: int   98 172 185 78 77 100 117 54 96 83 ...

> lm1 <- lm(phos~soil*topo,data=sp)
> leveneTest(lm1)

Levene's Test for Homogeneity of Variance (center = median)
      Df F value Pr(>F)
group  7  0.3741 0.9043
      16

> residPlot(lm1)
```



```
> adTest(lm1$residuals)

Anderson-Darling normality test

data:  lm1$residuals
A = 0.2126, p-value = 0.8351

> outlierTest(lm1)

No Studentized residuals with Bonferonni p < 0.05
Largest |rstudent|:
      rstudent unadjusted p-value Bonferonni p
1 -2.824098          0.012821      0.30769
```

```
> anova(lm1)
```

Analysis of Variance Table

Response: phos

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
soil	1	17876.0	17876.0	22.9818	0.0001988
topo	3	9693.8	3231.3	4.1542	0.0235128
soil:topo	3	11390.8	3796.9	4.8814	0.0134826
Residuals	16	12445.3	777.8		

When a two-way ANOVA model has a significant interaction term then multiple comparisons must be computed between each group that can be identified as combinations of the two factors. Unfortunately, this is not a straightforward calculation with the `glht()` function. However, the calculation can be made relatively easily by creating a single factor that consists of the combinations of the two original factors, fitting a one-way ANOVA model to this new single factor, and then submitting this result to the `glht()` function. This process is illustrated below.

```
> sp$comb <- sp$soil:sp$topo
> view(sp)
```

	soil	topo	phos	comb
7	shale	south	117	shale:south
11	shale	hilltop	12	shale:hilltop
13	sandstone	valley	19	sandstone:valley
16	sandstone	north	27	sandstone:north
17	sandstone	north	49	sandstone:north
21	sandstone	south	72	sandstone:south

```
> lm1a <- lm(phos~comb,data=sp)
> anova(lm1a)
```

Analysis of Variance Table

Response: phos

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
comb	7	38961	5565.8	7.1555	0.0005729
Residuals	16	12445	777.8		

```
> spint.mc <- glht(lm1a, mcp(comb="Tukey"))
> summary(spint.mc)
```

Simultaneous Tests for General Linear Hypotheses

Multiple Comparisons of Means: Tukey Contrasts

Fit: `lm(formula = phos ~ comb, data = sp)`

Linear Hypotheses:

	Estimate	Std. Error	t value	Pr(> t)
sandstone:north - sandstone:hilltop == 0	1.667	22.772	0.073	1.00000
sandstone:south - sandstone:hilltop == 0	19.333	22.772	0.849	0.98682
sandstone:valley - sandstone:hilltop == 0	-4.000	22.772	-0.176	1.00000
shale:hilltop - sandstone:hilltop == 0	4.667	22.772	0.205	1.00000
shale:north - sandstone:hilltop == 0	53.333	22.772	2.342	0.33035

```

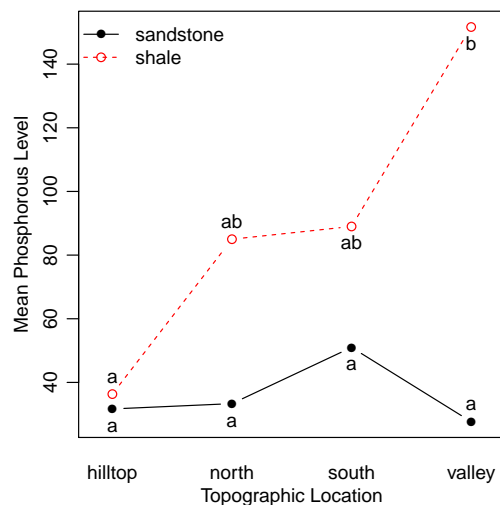
shale:south - sandstone:hilltop == 0      57.333      22.772      2.518      0.25579
shale:valley - sandstone:hilltop == 0     120.000      22.772      5.270      0.00145
sandstone:south - sandstone:north == 0     17.667      22.772      0.776      0.99218
sandstone:valley - sandstone:north == 0    -5.667      22.772     -0.249      1.00000
shale:hilltop - sandstone:north == 0        3.000      22.772      0.132      1.00000
shale:north - sandstone:north == 0         51.667      22.772      2.269      0.36548
shale:south - sandstone:north == 0         55.667      22.772      2.445      0.28527
shale:valley - sandstone:north == 0        118.333      22.772      5.196      0.00180
sandstone:valley - sandstone:south == 0    -23.333      22.772     -1.025      0.96345
shale:hilltop - sandstone:south == 0       -14.667      22.772     -0.644      0.99747
shale:north - sandstone:south == 0         34.000      22.772      1.493      0.80043
shale:south - sandstone:south == 0         38.000      22.772      1.669      0.70534
shale:valley - sandstone:south == 0        100.667      22.772      4.421      0.00788
shale:hilltop - sandstone:valley == 0        8.667      22.772      0.381      0.99992
shale:north - sandstone:valley == 0         57.333      22.772      2.518      0.25516
shale:south - sandstone:valley == 0         61.333      22.772      2.693      0.19393
shale:valley - sandstone:valley == 0       124.000      22.772      5.445      0.00109
shale:north - shale:hilltop == 0           48.667      22.772      2.137      0.43376
shale:south - shale:hilltop == 0           52.667      22.772      2.313      0.34385
shale:valley - shale:hilltop == 0          115.333      22.772      5.065      0.00215
shale:south - shale:north == 0              4.000      22.772      0.176      1.00000
shale:valley - shale:north == 0            66.667      22.772      2.928      0.13113
shale:valley - shale:south == 0            62.667      22.772      2.752      0.17669
(Adjusted p values reported -- single-step method)

```

```

> fitPlot(lm1,change.order=TRUE,interval=FALSE,main="",ylab="Mean Phosphorous Level",
  xlab="Topographic Location",legend="topleft")
> addSigLetters(lm1,change.order=TRUE,lets=c("a","a","a","ab","a","ab","a","b"),
  pos=c(1,3,1,3,1,1,3,1))

```



The following code can be used to isolate the multiple comparisons that have a p-value less than 0.05.

```

> spmc <- summary(spint.mc)
> names(spmc$test$coefficients)[spmc$test$pvalues<0.05]

```

```
[1] "shale:valley - sandstone:hilltop" "shale:valley - sandstone:north"  
[3] "shale:valley - sandstone:south"   "shale:valley - sandstone:valley"  
[5] "shale:valley - shale:hilltop"
```